

CertVIC V9 Scaffold: Pilot Evidence and Model-Dependent Specificity

Anonymous

The V9 scaffold frames CertVIC as a pilot-scale certification workflow with model-dependent specificity limitations. Main-500 is not claimed. Related work should position visual intervention certification, VLM robustness controls, and uncertainty-aware evaluation without overstating novelty. The method separates update-gap estimation, spurious specificity gates, quality checks, and human review. Experiments currently consist of the reviewed Main-200 pilot and V8/V8.1 diagnostics. Spurious V2 and Main-500 are prepared but deferred. Existing evidence shows InternVL and LLaVA pass the V1 irrelevant-edit specificity gate, while Qwen fails it at 12/94. This supports model-dependent specificity only.

Specificity controls. We evaluate irrelevant-edit specificity as a separate gate from the main update-gap result. InternVL and LLaVA pass the existing irrelevant-edit specificity gate, while Qwen2.5-VL-7B fails it under the V1 control set with 12 flips out of 94 items. A stricter V9 Spurious V2 control set has been prepared for a preregistered rerun, but those GPU results are not yet available. Until that run is imported and validated, the paper must not claim all-model specificity.

Model-dependent specificity limitation. The present evidence supports model-dependent specificity, not all-model specificity. In particular, Qwen shows a strong update-gap result but elevated sensitivity to irrelevant controls. Main conclusions must therefore be conditioned on this limitation: the control evidence is strongest for InternVL and LLaVA, while Qwen requires either a successful stricter Spurious V2 rerun or an explicit model-specific limitation.

Inference plan. The primary estimand is the paired answer-change/update gap on reviewed items. Spurious specificity is a separate gate with threshold ≤ 0.10 and is not weakened after observing failures. Sequential reports must use anytime-valid confidence sequences and must keep exclusion sensitivity analyses separate from canonical results.

Diagnostic mechanism and polarity probes. Polarity and mechanism probes are treated as diagnostic-only. They help identify answer-stability and prompt-sensitivity patterns, but they are not certification evidence and are not used to relax the specificity threshold.

Qualitative findings. The V9 gallery uses deterministic inclusion rules, including all twelve Qwen spurious-control failures. The gallery supports failure analysis and limitations, not standalone certification claims. Main limitations are missing Spurious V2 outputs, pending real human review for Qwen failures, absent Main-500 results, and dataset/domain narrowness. CertVIC is promising but not CVPR-main-ready without the deferred gates.